

Probabilistic attacks against compressed encrypted protocols

Dimitrios Karakostas

Dionysis Zindros

Aristeidis Pagourtzis

Theoretical background

(gzip)

- **gzip**: The most used encryption software in the Internet.
- Implements the DEFLATE algorithm:

$$DEFLATE(m) = Huffman(LZ77(m))$$

Theoretical background

(LZ77)

- **LZ77**: Lossless data compression algorithm, published in 1977 by A. Lempel and J. Ziv.
- Method:
 - Find repeated portions of data.
 - Replace them with references as [length, offset].
 - Minimum length = 3.
 - Maximum offset = 32Kb.

Theoretical background

(LZ77)


LZ77 example

Hello, world! I love you.

Hello, world! I hate you.

Hello, world! Hello world! Hello world!

Hello, world! I love you.



Theoretical background

(Huffman)

- **Huffman coding:** Lossless data compression algorithm, proposed by D. Huffman in 1952.
- Method:
 - Analyze the frequency of each letter in the text.
 - Replace common letters with short codes.
 - Replace rare letters with long codes.
 - Code alphabet should be prefix free.

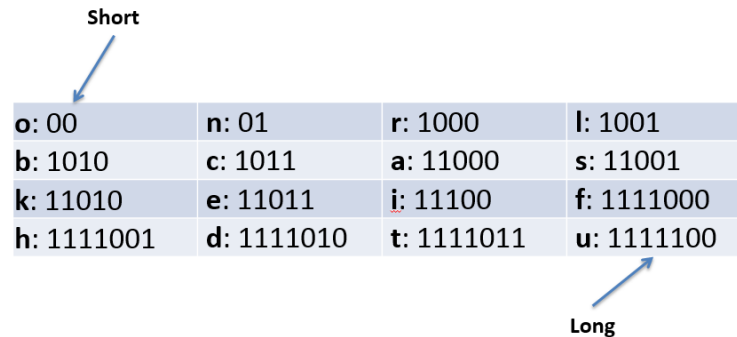
Theoretical background (Huffman)

Huffman example

- Frequency analysis:

o: 6	n: 5	r: 3	l: 3
b: 3	c: 3	a: 3	s: 2
k: 2	e: 2	i: 2	f: 2
h: 1	d: 1	t: 1	u: 1

- Code alphabet:



o: 00	n: 01	r: 1000	l: 1001
b: 1010	c: 1011	a: 11000	s: 11001
k: 11010	e: 11011	i: 11100	f: 1111000
h: 1111001	d: 1111010	t: 1111011	u: 1111100

Theoretical background

(Same-origin policy)

- **Same-origin policy:** scripts in one page are allowed to access data in a second page if both have the same origin.
- Origin: protocol, host and port of a URL.
- Documents retrieved from distinct origins are isolated from each other.
- i.e. a document retrieved from <http://example.com/target.html> is disallowed to access the DOM of a document retrieved from <https://head.example.com/target.html>.

Theoretical background

(Same-origin policy)

- Attacks on same-origin policy:
 - Cross-site scripting (XSS): vulnerability that allows an attacker to inject a client-side script into web pages viewed by other users.
 - Cross-site request forgery (CSRF): exploit that allows the attacker to issue unauthorized requests to a website, on behalf of a user the website trusts.

Theoretical background

(TLS)

- **Transport Layer Security (TLS):** protocol that provides security over the internet.
- Prevents eavesdropping, tampering or message forgery.
- TLS handshake allows the negotiation of a symmetric key via asymmetric cryptography, provided by certificates created by trusted authorities.

Theoretical background (TLS)

- TLS record structure

+	Byte +0	Byte +1	Byte +2	Byte +3
Byte 0	Content type			
Bytes 1..4	Version		Length	
	(Major)	(Minor)	(bits 15..8)	(bits 7..0)
Bytes 5.. $(m-1)$	Protocol message(s)			
Bytes $m..(p-1)$	MAC (optional)			
Bytes $p..(q-1)$	Padding (block ciphers only)			

Theoretical background

(MitM)

- Man-in-the-Middle: one of the most common attack vectors on modern communications.



- Common MitM techniques:
 - ARP Spoofing: the attacker sends ARP messages, so that its MAC address is associated with the target endpoint's IP address.
 - DNS Poisoning: the attacker introduces data into a DNS resolver's cache, to return incorrect address for the chosen endpoint.

IND-PCPA

(PCPA game)

- Traditionally, cryptographers have used games for security analysis
- IND-CPA, IND-CCA{1,2}
- We introduce a new security game:

Indistinguishability under partially chosen plaintext attack (IND-PCPA)

IND-PCPA

(PCPA game)

- The challenger generates a pair P_k, S_k and publishes P_k to the adversary.
- The adversary may perform a polynomially bounded number of encryptions or other operations.
- Eventually, the adversary submits two distinct chosen plaintexts M_0, M_1 to the challenger.
- The challenger selects a bit $b \in \{0,1\}$ uniformly at random.
- The adversary can then submit any number of selected plaintexts $R_i, i \in N, |R| \geq 0$, and the challenger sends the ciphertext $C_i = E(P_k, M_b | R_i)$ back to the adversary.
- The adversary is free to perform any number of additional computations or encryptions, before finally guessing the value of b .

IND-PCPA

(PCPA game)

A cryptosystem is indistinguishable under partially chosen plaintext attack, if every probabilistic polynomial time adversary has only a negligible advantage on finding b over random guessing.

IND-PCPA

- IND-PCPA vs IND-CPA:
 - The adversary submits the empty string as chosen plaintext.
 - The challenger then sends back:
$$C = E(P_k, M_b \parallel "") = E(P_k, M_b)$$
 - which is the challenger response of the IND-CPA game.
 - Intuitively, if the adversary can beat the game of IND-PCPA, he also has the ability to beat IND-CPA.

IND-PCPA

- PCPA scenario on compression-before-encryption protocol:
 - A system creates:
 $c = \text{Encrypt}(\text{Compress}(m))$
where c is the ciphertext of the compressed m .
 - The attacker issues a PCPA creating:
 $m = n_1 || \text{secret} || n_2 || \text{reflection} || n_3$
 - where n_1, n_2, n_3 are random nonces.
 - If the chosen reflection is the same as the secret, a pattern emerges and the compression is better, possibly resulting in smaller ciphertext, compared to the one of a wrong reflection.

IND-PCPA

(PCPA exploits)

- CRIME:
 - [Rizzo, Duong '12]
 - CRIME attacked TLS header compression in HTTPS.
 - TLS header compression is now disabled.
 - CRIME is no longer possible.
 - CRIME set the foundation for compression/encryption attacks.

IND-PCPA

(PCPA exploits)

- BREACH:
 - [Prado, Harris, Gluck '13]
 - BREACH was based on CRIME.
 - BREACH attacks HTTPS response.
 - Original BREACH attack had specific assumptions:
 - Against stream ciphers.
 - No noise in response.
 - Secret has known prefix, bootstrapping is trivial.

Attack model

(Assumptions)

- The attacker has gained control of the victim's network and can view the victim's encrypted traffic, which can be accomplished by MitM.
- The attack script issues requests toward the chosen endpoint from the victim's browser, i.e. via XSS.
- Each request contains a chosen stream of data, which is reflected in the response body, along with the secret.
- Compression is applied on both the secret and the reflection.

Attack model

- MitM implementation:
 - We add a rule in the hosts file of the lab machine, in order for all traffic toward an endpoint to be redirected to the localhost.
 - We implemented a Python MitM proxy, that opens TCP sockets on both the lab machine and the endpoint and forwards traffic on both ends, while parsing the header and (encrypted) body TLS record.
 - We also implemented a defragmentation mechanism, in order to parse records that span over multiple TCP packets.

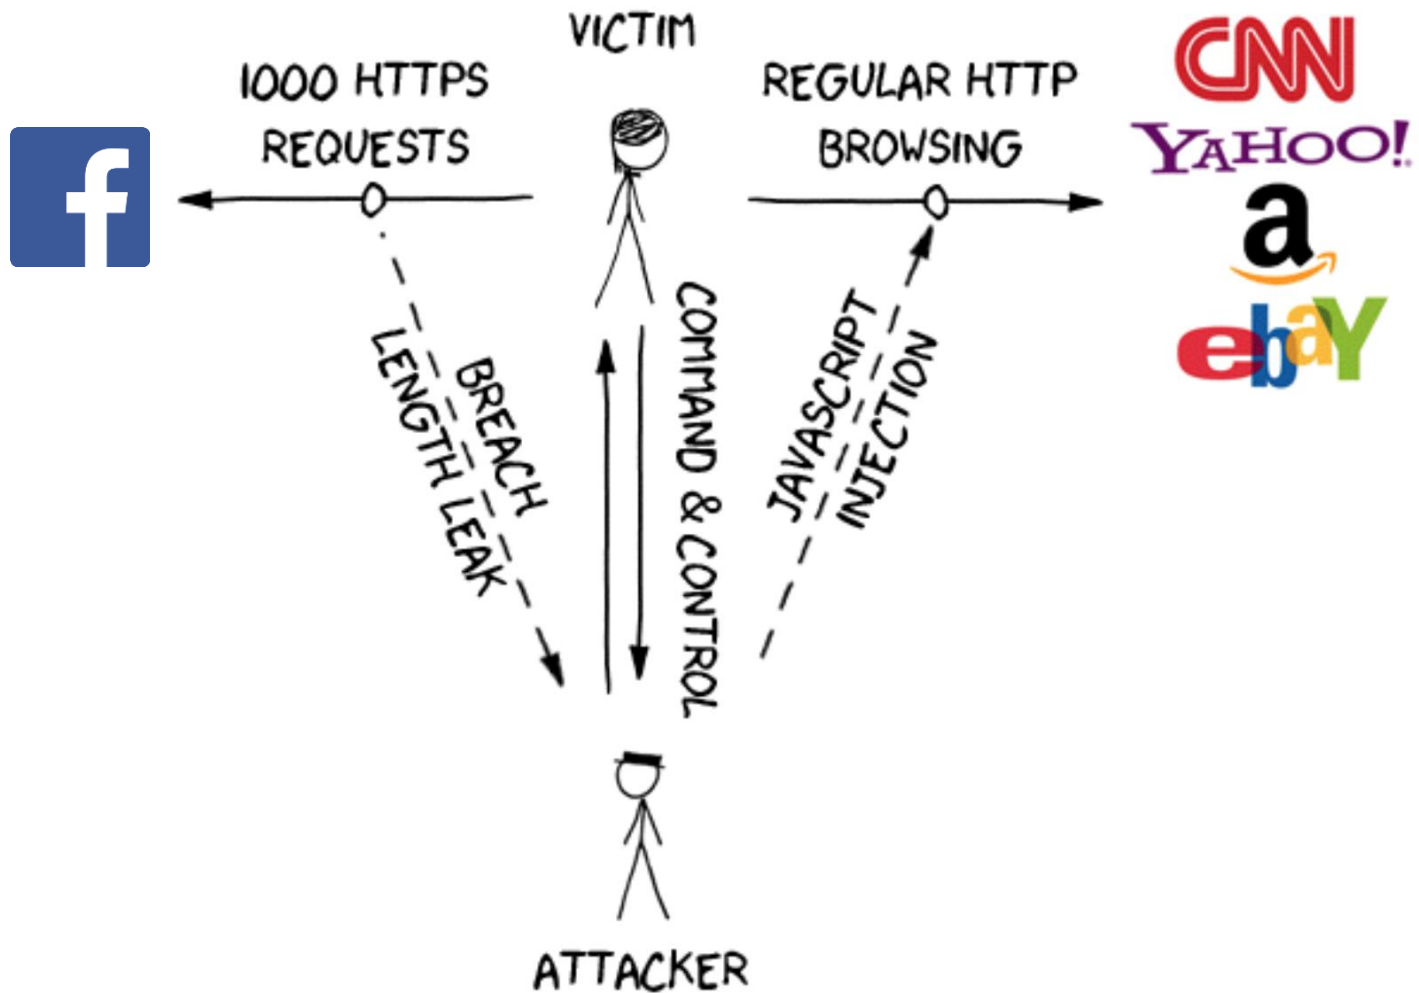
Attack model

- BREACH script implementation:
 - The user inputs a known prefix for the secret, needed to bootstrap the attack, and the alphabet that the characters of the secret belong.
 - An attack vector is created, with each item corresponding to a fragment of the alphabet, where the sum of the fragment makes up the whole alphabet.
 - A request is issued for each item of the vector every 4 seconds, resuming from the beginning when the end of the vector is reached.
 - The requests are made in the form of tags, injected in the HTML body of a controlled website.

Attack model

- Attack persistence:
 - We propose a command-and-control mechanism that allows the execution of the attack without the need of a contaminated website, that the victim would visit.
 - The victim needs to browse the HTTP web.
 - The attacker that controls the victim's traffic would inject the attack script in the response from a regular HTTP website.

Attack model



Attack model

- Vulnerable endpoints:
 - Facebook Chat messages
 - Gmail Authentication token
 - Gmail private emails

Attack model

(Facebook Chat messages)

- Facebook Chat messages:
 - Facebook provides a lightweight mobile version, Facebook Touch.
 - It also allows a search functionality via URL, in the form:
https://touch.facebook.com/messages?q=<search_string>
 - The search string is reflected in the body of the response.
 - Also, regardless of the search results, the last message of the 5 most recent conversations is also included in the body.

Attack model

(Facebook Chat messages)

Noise

[illegible]

- Private message

[illegible]

Reflection

Attack model

(Gmail Authentication token)

- Gmail Authentication token:
 - Gmail provides a plain HTML version for faster browsing, which enables a search functionality as:
https://mail.google.com/mail/u/0/x/?s=q&q=<search_string>
 - Each request should contain a valid, random-generated string between the *0* and *x* parameter of the URL.
 - If no string is included, a redirection to a URL that contains such a string is applied, returning an empty result page, stating the action as incomplete.
 - However, the HTML body contains both the search string and the authentication token for the account.
 - Different tokens of different accounts demonstrate a fixed prefix: “AF6bup”.

Attack model

(Gmail Authentication token)

```
amp;pv=tl&amp;eot=1&
amp;q=ladbfsk!1_3_5_7_9_b_d_f_h_j_l_n_p_r_t_v_x_z_B_D_F_H_J_L_N_P_R_T_V_X_Z__2_4_
6_8_a_c_e_g_i_k_m_o_q_s_u_w_y_A_C_E_G_I_K_M_O_Q_S_U_W_Y_-_AF6bup0znq&v=b&
amp;s=q">Compose</a></td></tr></table><div class="notification">We cannot
complete the action at this time. Please try again using the search action above.
</div><form action="?&mnut=tl&v=mnu" name="f" method="post"><input
type="hidden" name="at" value="AF6bupNx9G8BD_Wr7frvMfpnj_j_Nh_0GVQ" /><input
type="hidden" name="nredir" value="?&at=AF6bupNx9G8BD_Wr7frvMfpnj_j_Nh_0GVQ&
amp;s=q" /><input type="hidden" name="nredir" value="?&
amp;q=ladbfsk!1_3_5_7_9_b_d_f_h_j_l_n_p_r_t_v_x_z_B_D_F_H_J_L_N_P_R_T_V_X_Z__2_4_
6_8_a_c_e_g_i_k_m_o_q_s_u_w_y_A_C_E_G_I_K_M_O_Q_S_U_W_Y_-_AF6bup0znq&s=q"
/><input type="hidden" name="search" value="query" /><div class="noMatches">No
results for.
ladbfsk!1_3_5_7_9_b_d_f_h_j_l_n_p_r_t_v_x_z_B_D_F_H_J_L_N_P_R_T_V_X_Z__2_4_6_8_a
_c_e_g_i_k_m_o_q_s_u_w_y_A_C_E_G_I_K_M_O_Q_S_U_W_Y_-_AF6bup0znq</div><script
type="text/javascript">
var token="AF6bupNx9G8BD_Wr7frvMfpnj_j_Nh_0GVQ";var
searchPageLinks=document.getElementsByClassName("searchPageLink");
for(i=0;i<searchPageLinks.length;i++)searchPageLinks[i].onclick=function(e){var
href=e.currentTarget.href;var form=document.createElement("form");
form.setAttribute("method","post");form.setAttribute("action",href);var
inputToken=document.createElement("input");
```

Reflection

Authentication token

Attack model

(Gmail private emails)

- Gmail private emails:
 - The attacker issues a search request through a URL like:
https://mail.google.com/mail/u/0#search/<search_string>
 - The response body does not include the search string, however, it contains both the Subject and a fragment of the body of the latest inbox mails.
 - The attacker could send multiple mails to the victim, that would be included in the response, along with other private mails.

Attack model

- Validation of secret-reflection compression:
 - We use mitmproxy¹, to extract the compressed body of a response that was obtained with the attack.
 - We use infgen², to disassemble the compressed body to the LZ77 compression of the initial data stream.

```
match 21 378
match 8 325
literal 'potrymakwi_a_b_c_d_e_f_g_h_i_j_k_l_m_n_o_p_q_r_s_t_u_v_w_x_y_z'
match 4 19121
literal 'j'
match 3 79
literal 'p'
match 4 24700
match 6 16724
match 3 85
match 6 3090
match 8 1330
```

Bad compression

```
match 21 378
match 8 325
literal 'potrymakwi_a_b_c_d_e_f_g_h_i_j_k_l_m_n_o_p_q_r_s_t_u_v_w_x_y_z'
match 5 19124
match 3 79
literal 'p'
match 4 24724
match 6 16727
match 3 85
match 6 3090
match 8 1330
literal 's...
```

Good compression

1. <https://mitmproxy.org>
2. <http://www.zlib.net/infgen.c.gz>

Statistical methods

(Block ciphers)

- Original attacks assumed stream ciphers. e.g. original BREACH assumed RC4.
- [Prado, Neal, Gluck] suggested block ciphers are vulnerable, but did not provide practical attack details.
- In this work, we perform practical attacks against popular block ciphers:
 - We attack AES_128 used in Facebook, Gmail, Twitter, Wikipedia, YouTube, Amazon etc.
- We have found that the AES implementation in the NSS library displays certain patterns.

Statistical methods

(Block ciphers)

```
User application payload: 1083
Endpoint application payload: 40
Endpoint application payload: 1524
Endpoint application payload: 101
Endpoint application payload: 1524
Endpoint application payload: 1104
Endpoint application payload: 1524
Endpoint application payload: 2604
Endpoint application payload: 1351
User application payload: 40
User application payload: 1083
Endpoint application payload: 40
Endpoint application payload: 1524
Endpoint application payload: 101
Endpoint application payload: 1524
Endpoint application payload: 1104
Endpoint application payload: 1524
Endpoint application payload: 2604
Endpoint application payload: 1353
User application payload: 40
```

} First request

} Second request

Facebook flow

```
User application payload: 255
Endpoint application payload: 270
Endpoint application payload: 350
Endpoint application payload: 41
User application payload: 41
User application payload: 259
Endpoint application payload: 74
Endpoint application payload: 1395
Endpoint application payload: 1287
Endpoint application payload: 41
User application payload: 41
User application payload: 255
Endpoint application payload: 271
Endpoint application payload: 402
Endpoint application payload: 41
User application payload: 41
User application payload: 260
Endpoint application payload: 70
Endpoint application payload: 1395
Endpoint application payload: 1304
Endpoint application payload: 41
User application payload: 41
```

} First request

} First redirection

} Second request

} Second redirection

Gmail flow

Statistical methods

(Block ciphers)

```
User application payload: 3142
Endpoint application payload: 214
Endpoint application payload: 340
Endpoint application payload: 36
User application payload: 3161
User application payload: 36
Endpoint application payload: 78
Endpoint application payload: 229
Endpoint application payload: 36
User application payload: 36
User application payload: 3015
Endpoint application payload: 53
Endpoint application payload: 1122
Endpoint application payload: 36
User application payload: 36
```

```
User application payload: 3142
Endpoint application payload: 80
Endpoint application payload: 340
Endpoint application payload: 36
User application payload: 36
User application payload: 3160
Endpoint application payload: 67
Endpoint application payload: 230
Endpoint application payload: 36
User application payload: 36
User application payload: 3015
Endpoint application payload: 53
Endpoint application payload: 1125
Endpoint application payload: 36
User application payload: 36
```

Old browser flow

```
User application payload: 2220
Endpoint application payload: 98
Endpoint application payload: 362
Endpoint application payload: 41
User application payload: 41
User application payload: 2105
Endpoint application payload: 46
Endpoint application payload: 1330
Endpoint application payload: 41
User application payload: 41
User application payload: 2205
Endpoint application payload: 237
Endpoint application payload: 418
Endpoint application payload: 41
```

```
User application payload: 2220
User application payload: 41
Endpoint application payload: 98
Endpoint application payload: 259
Endpoint application payload: 41
User application payload: 41
User application payload: 2105
Endpoint application payload: 46
Endpoint application payload: 1306
Endpoint application payload: 41
User application payload: 41
User application payload: 2205
Endpoint application payload: 236
Endpoint application payload: 424
Endpoint application payload: 41
User application payload: 41
```

Newer browser flow

Statistical methods

(Block ciphers)

- We issue a large amount of requests for each item of the attack vector.
- We calculate the mean response length for each item.
- The correct guess should converge to smaller mean response length, compared to the others.

Statistical methods

(Huffman fixed-point)

- Huffman tables may be tampered, when different requests are issued.
- We describe a methodology to bypass this Huffman-induced noise:
 - An alphabet pool is created, containing every item in the alphabet of the secret.
 - In each request, the part of the alphabet that is not being tested is appended in the beginning.
 - Each request presents same letter frequency, although the text is rearranged.

Statistical methods

(Huffman fixed-point)

?q=rynmkwi_1_2_3_4_5_6_7_8_9_Credit Card: 0znq
?q=rynmkwi_0_2_3_4_5_6_7_8_9_Credit Card: 1znq
?q=rynmkwi_0_1_3_4_5_6_7_8_9_Credit Card: 2znq
?q=rynmkwi_0_1_2_4_5_6_7_8_9_Credit Card: 3znq
?q=rynmkwi_0_1_2_3_5_6_7_8_9_Credit Card: 4znq
?q=rynmkwi_0_1_2_3_4_6_7_8_9_Credit Card: 5znq
?q=rynmkwi_0_1_2_3_4_5_7_8_9_Credit Card: 6znq
?q=rynmkwi_0_1_2_3_4_5_6_8_9_Credit Card: 7znq
?q=rynmkwi_0_1_2_3_4_5_6_7_9_Credit Card: 8znq
?q=rynmkwi_0_1_2_3_4_5_6_7_8_Credit Card: 9znq

Statistical methods

(Hill-climbing parallelization)

- The alphabet partitioning follows a divide-and-conquer scheme.
- Example:
 - The attack vector on digits could be as follows:
[“0 2 4 6 8”, “1 3 5 7 9”]
 - The correct digit will be compressed with the secret, so the vector item that contains it will present better behavior.
 - Each stage of the attack outputs a chosen half of the tested alphabet fragment, until the chosen half contains only one digit, which is the correct one.
- This method could reduce the time of the attack from $O(|S|)$ to $O(\log |S|)$.

Statistical methods

(Cross-domain parallelization)

- Most websites use subdomains for specific applications, such as mobile versions.
- Cookies from the parent domain are available to the subdomains.
- If the subdomains handle similar data, containing the chosen secret, the attack could be issued against them.
- The parallelization could effectively increase the attack efficiency up to Nx , where N is the number of different subdomains.

Statistical methods

(Point-system meta-predictor)

- Experiments revealed that the correct guess does not always result in minimum mean response length.
- However, the correct item is more probable to be among the *best* ones over time, compared to the others, that may demonstrate only a spike in performance for a certain period.
- For that reason we introduce a point-system that evaluates the performance of each item compared to the others.

1: 20	2: 16
3: 12	4: 10
5: 8	6: 6
7: 4	8: 3
9: 2	10: 1

Experimental results


(Facebook Chat messages)

- We created a lab account, that has no friends, no user activity of any kind, except for a self-sent private message, containing the secret.
- We choose a prefix to bootstrap the attack, while the alphabet consists of lowercase and uppercase letters.
- We issue the attack using the serial method of requests, performing 4000 iterations, with a 4 second interval between requests.


Total time

$$4000 * 52 * 4 = 832000 \text{ seconds} = 9 \text{ days}$$

Experimental results (Facebook Chat messages)



Experimental results (Facebook Chat messages)



Experimental results

(Gmail Authentication token)


- We use the hill-climbing parallelized attack method to steal the auth token of a regular Gmail account.
- The alphabet consists of lowercase, uppercase, digits and dashes, so the stages of the attack are $\log(64) = 6$.
- We repeat each stage of the attack, until one of the two halves is chosen 4 times, so at most 7 attempts are made for each stage of the parallelization.

Total time

$$4000 * 7 * 6 * 4 = 672000 \text{ seconds} = 7 \text{ days}$$

Experimental results

(Gmail Authentication token)



Mitigation techniques

- [Prado etc.] proposed several mitigation techniques:
 - **Length hiding.** In this work, we were able to defeat this mitigation measure through noise by-passing.
 - **Separating secrets from user input.** In this work, we were able to defeat this mitigation measure through alternative secrets: Secrets and user input are sometimes one and the same, e.g. private messages.
 - **Masking secrets.** This mitigation mechanism is still feasible. But we showed that many more secrets than CSRF tokens must be masked.
 - **Rate limiting and monitoring.** This mitigation mechanism is still feasible.
 - **CSRF protection.** In this work, we showed that this is not adequate mitigation, as secrets other than CSRF can be stolen.
 - **Disabling compression.** While this solves the problem, it is not a practical solution.

Novel mitigation techniques

(Compressibility annotation)

- We propose that web servers and web application servers cooperate to indicate which portions must not be compressed.
- Web application server returns annotated response:
 - Annotation indicates where secrets are located.
 - Annotation indicates where reflection is located.
 - Annotation uses some special format.
- Must be implemented separately in every web framework, e.g. Django, Ruby on Rails.
- Web server interprets annotated web application server response and changes compression behavior.
- Annotated reflections and secrets always sent as literals
- Must be implemented separately in web servers, e.g. mod_breach for Apache, Nginx etc.

Novel mitigation techniques (SOS headers)

- [Schema, Toukharian '13] propose SOS headers as an extension to CSP.
- A policy applies to each cookie, specifying whether it should be included in a request.
- Policies applied: any, self, isolate
- Pre-flight requests are made to check for exceptions.
- If trusted websites use HSTS policy and cookies are not included in other cases, the response would not contain the secret.
- Complete mitigation of the attack.

Conclusion

- Our contributions:
 - Definition of IND-PCPA
 - Attack optimization:
 - Parallelization
 - Point-system prediction
 - Attack persistence
 - Alternative secrets
 - Experimental results on major systems
- Future work:
 - Mathematical proof for IND-PCPA properties
 - HTTP injection persistency mechanism
 - Integration of MitM attacks
 - Implementation of proxy on TCP level
 - Implementation of novel mitigation techniques

Thank you!

Questions?